

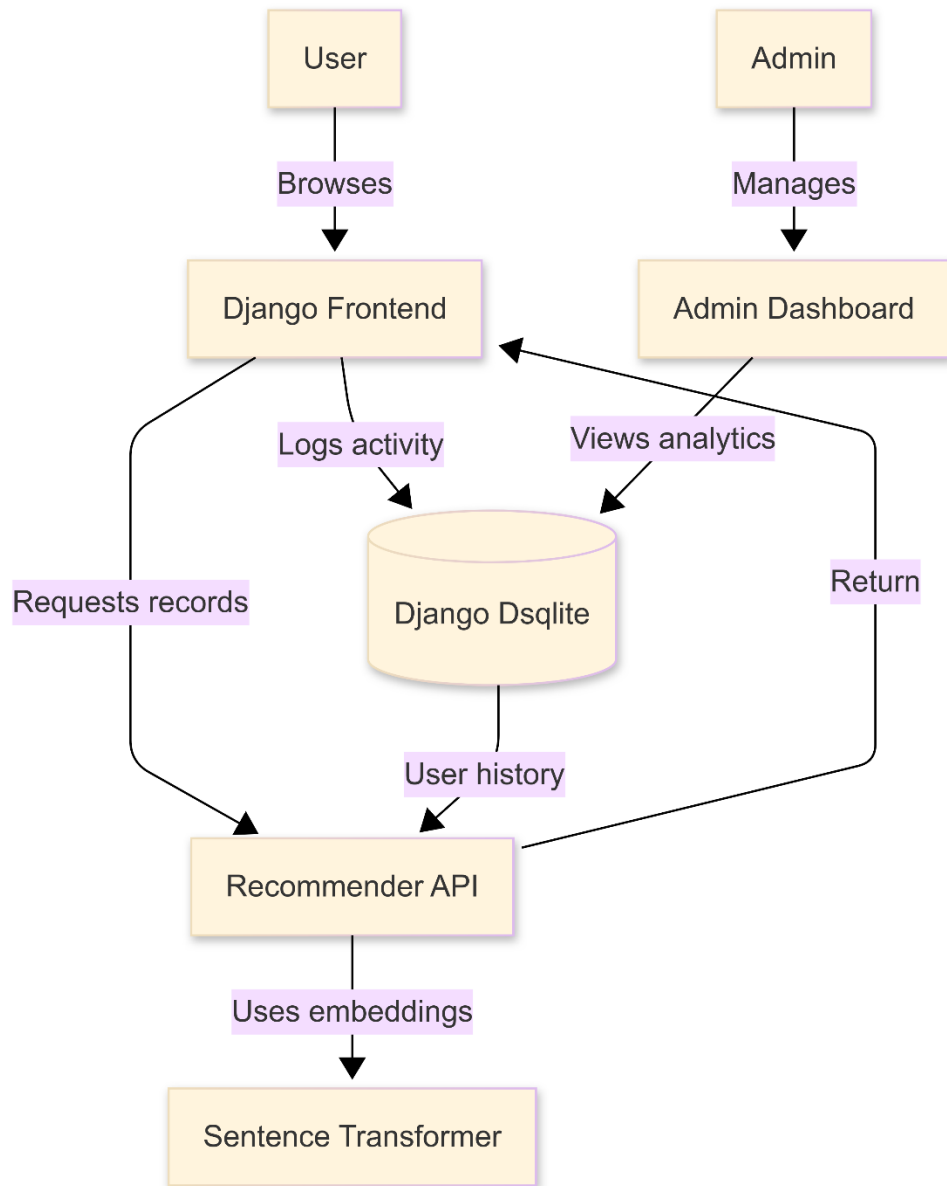
# E-Commerce Recommendation System

## Project Documentation

### 1. Introduction

A hybrid recommendation system combining:

- Django** (Frontend + Backend)
- FastAPI** (Recommendation Engine)
- TF-IDF Model + Cosine Similarity** (Content-based filtering)



*Fig 1: System Architecture Diagram*

## 2. Models Used

### a. Core Recommendation Models

Model	Type	Description
TF-IDF Vectorizer	NLP Feature Extraction	Converts product text (name, description) to numerical vectors
Cosine Similarity	Similarity Metric	Measures similarity between product vectors
Rule-Based Filtering	Collaborative	Recommends products from categories the user has interacted with

### How They Work Together:

- a. TF-IDF analyzes product descriptions
- b. Cosine similarity finds similar products
- c. User activity data adds personalization

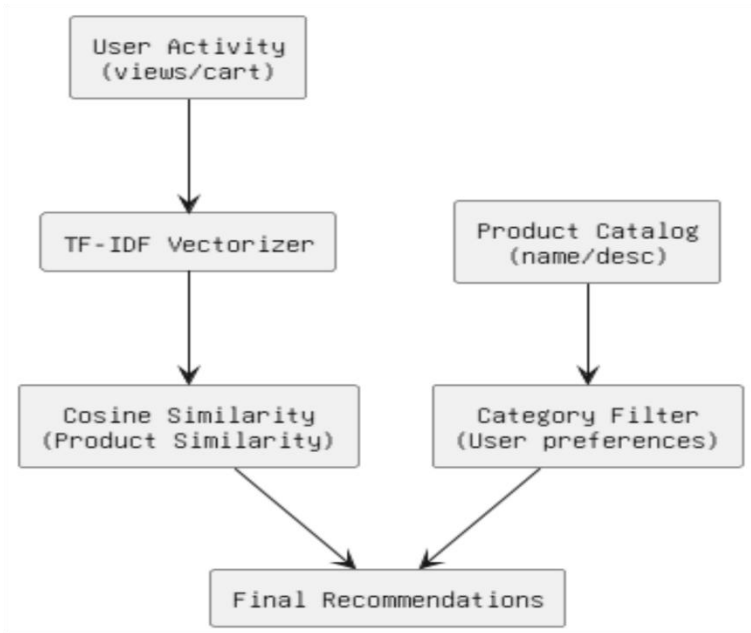
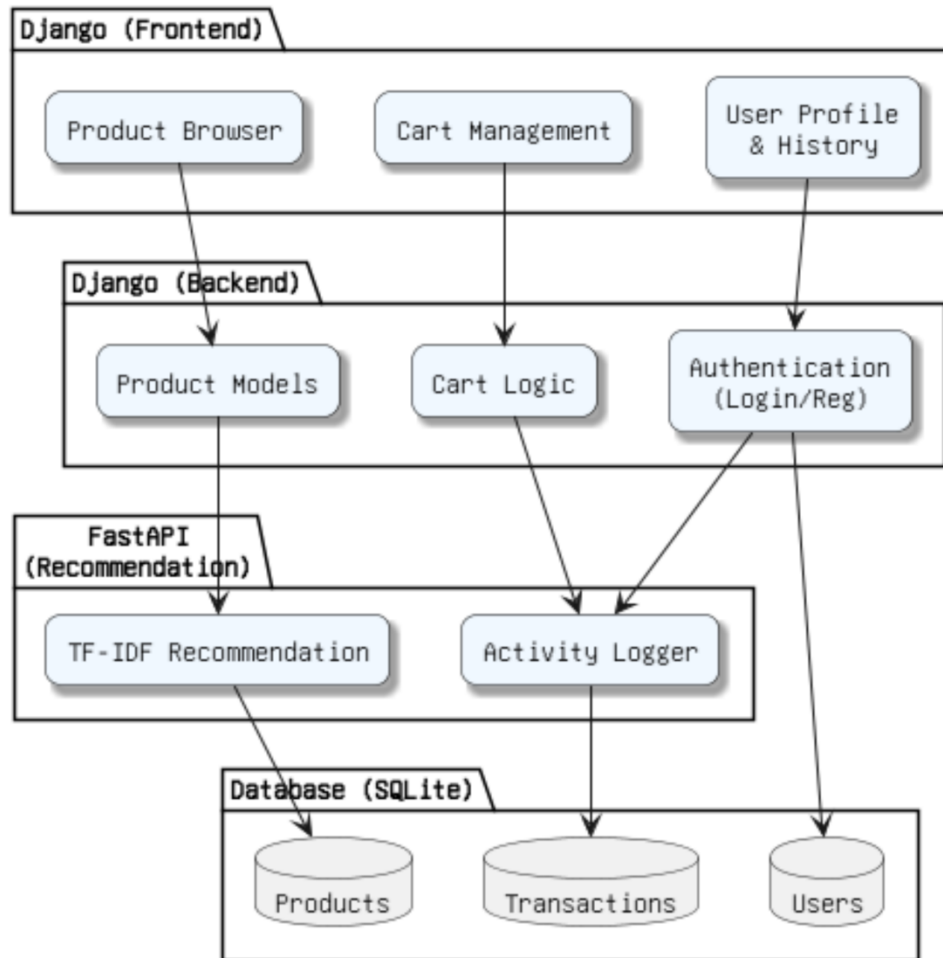


Fig 2: Recommendation Generation Process

### 3. System Architecture

#### a. Component Diagram



#### Key Parts :

- Django Frontend:** User interfaces
- Django Backend:** Business logic
- FastAPI:** Recommendation service
- Database:** Stores products/users/transactions in db.sqlite

#### Data Flow :

- User interacts with Django
- Activity sent to FastAPI
- Recommendations generated
- Displayed in Django templates

## 4. Implementation

### a. Django Components

Example model (products/models.py)

```
class Product(models.Model):  
    name = models.CharField(max_length=255)  
    description = models.TextField()  
    category = models.ForeignKey(Category, on_delete=models.CASCADE)
```

### b. FastAPI Endpoints

```
@app.get("/api/recommend")  
  
async def get_recommendations(user_id: int):  
    i. Get user activity  
    ii. Generate recommendations  
    iii. Return product IDs
```

## 5. Setup Guide

### a. Installation

pip install django fastapi scikit-learn

### Running the System

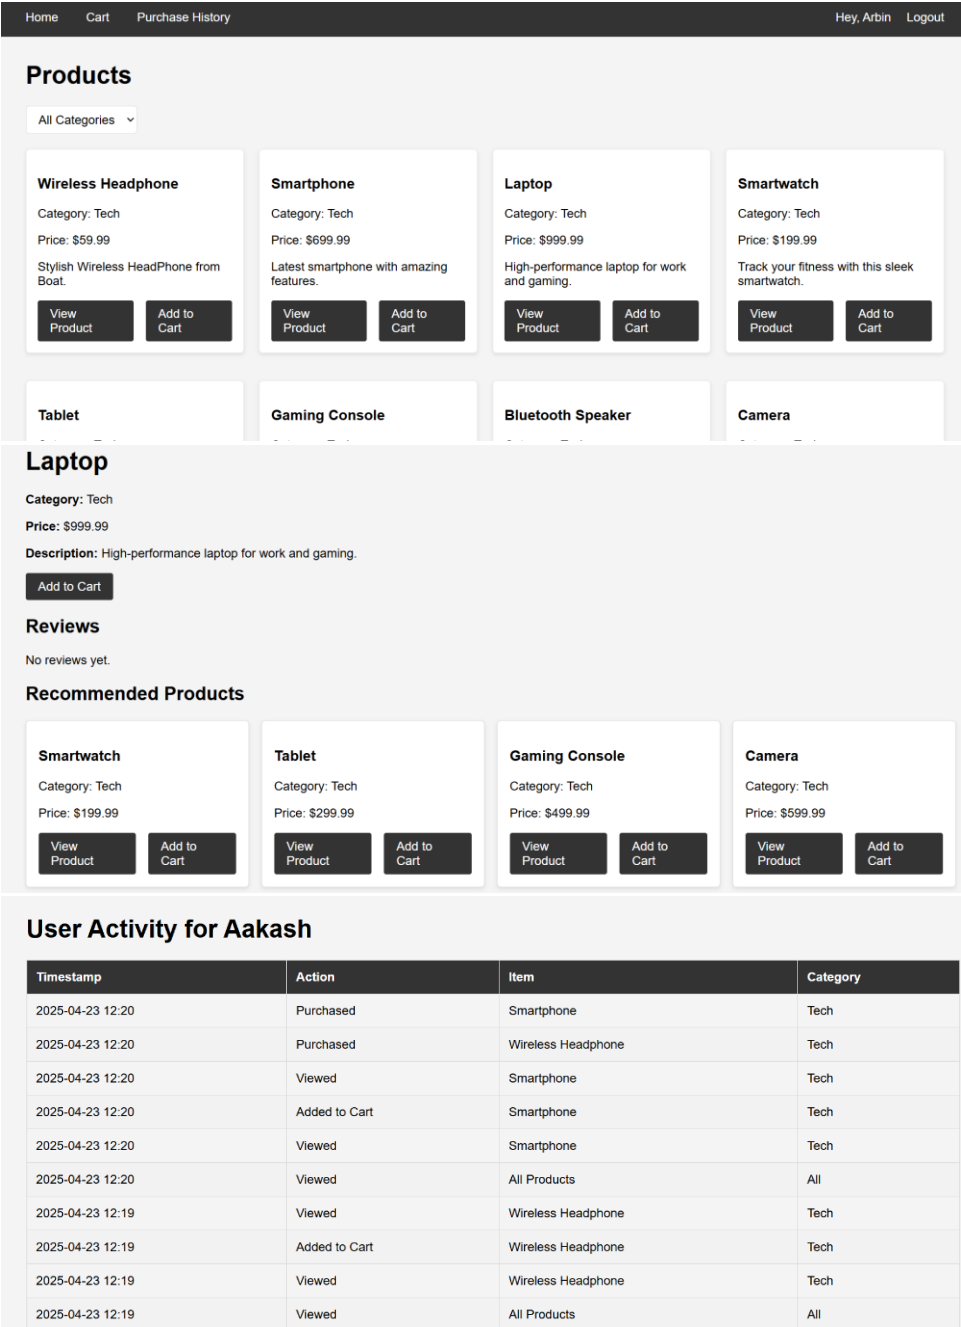
#### FastAPI (Port 8000)

uvicorn main:app --reload

#### Django (Port 8001)

python manage.py runserver 8001

## 6. Screenshots



## **7. Performance Notes**

**For 50 products(currently):**

- First-run latency: ~2s (generates all embeddings)
- Subsequent requests: ~50ms (cached)
- Memory usage: ~150MB (model + embeddings)

## **8. Limitations & Future Work**

**Current Limitations:**

- Basic text analysis only
- No real-time model updates

**Future Improvements:**

- Add BERT for semantic understanding
- Implement Redis caching